

IEOR 8100

Reinforcement Learning

Introductory presentation

Course Staff

- Instructor: Shipra Agrawal
 - Office hours: Wednesday 3:00pm-4:00pm Mudd 423
(starting next week)
- TA
 - Robin (Yunhao) Tang
PhD candidate, IEOR
 - Office hours TBD

ieor8100.github.io/rl/

Communication: *Piazza*

- No emails (unless absolutely necessary)
- Post questions on Piazza
 - Sign up for piazza
 - Can post publicly/privately/anonymously
- Announcements will be made on Piazza

Course requirements

- 4 lab assignments
- One paper presentation
- One research project

Research papers, project ideas and presentation schedule will be posted in two weeks.

Course Introduction

Reinforcement Learning

- Agent interacts and learns from a stochastic environment
- Science of sequential decision making
- Many faces of reinforcement learning
 - Reward systems (Neuro-science)
 - Classical/Operant Conditioning (Psychology)
 - Optimal control (Engineering)
 - Dynamic Programming (Operations Research)

Characteristics of Reinforcement Learning

- Sequential/online decisions
- No supervisor, only reward *signals*
- Feedback is delayed
- Actions effect observations (non i.i.d. training examples)


Examples

- Automated vehicle control/robotics
 - An unmanned helicopter learning to fly and perform stunts
- Game playing
 - Playing backgammon, Atari breakout, Tetris, Tic Tac Toe
- Medical treatment planning
 - Planning a sequence of treatments based on the effect of past treatments
- Chat bots
 - Agent figuring out how to make a conversation
 - Dialogue generation, natural language processing

Modeling foundation: MDP

- Markov Decision process: model for **sequential decisions**
 - Past information is captured by **state**
 - Agent takes an action, observes **new state and reward** generated from a stochastic model
 - Objective is some aggregate function of the individual rewards

Sequential decisions
Reward signals
(partial labels)
Delayed feedback
Actions effect observations



Reinforcement learning

- Reinforcement learning \equiv MDP with unknown stochastic model
- Agent observes samples : rewards, state transition
- Learn a good strategy (policy) for the MDP
 - Implicitly or explicitly learn the model dynamically from observations

The algorithm design problem

Design a strategy for taking actions sequentially, after observing current state

- Generate good reward
- Generate informative sample observations

and converge to optimal strategy

Challenges:

- Complex combination of learning and optimization
- There may be a tradeoff between reward and information
- Scale: large number of states, need to use structure

Course Goals

- Rigorous understanding of the MDP foundation:
 - Stochastic structure, algorithm design, convergence
- Conceptual understanding of recent algorithms for reinforcement learning
 - Mathematical insights into design principles
 - Occasional convergence results
- Ability to implement RL algorithms using some popular software platforms and simulators
 - Utilize Deep learning with tensorflow
 - OpenAI gym
- Ability to understand recent research papers
- New research!


Topics (tentative)

- **Introduction to MDP:** value-iteration, policy iteration, Q-value-iteration
- **Q-learning:** Tabular, function approximation
- **Deep Q-networks:** architecture, backpropagation, experience replay
- **Policy gradient methods:** Function approximation, Natural policy gradient, Trust region policy optimization, Actor critic methods,
- **Model-based RL**
- Further challenges: Exploration vs. exploitation, Adversarial training, Generalization, Multi-agent RL



Open AI gym

<https://gym.openai.com/envs/>

- Simulated environments for testing reinforcement learning algorithms



manoid-v1



Breakout-v0



Assignments: *Instabase*


- All software is pre-installed (Python+Jupyter+tensorflow+openAI gym)
- Create account using **username as UNI**, email as UNI@columbia.edu
 - Signup instructions <https://ieor8100.github.io/rl/cloudPlatform.html>
 - Signup requires a token.
- Access the lab assignments (Jupyter notebooks with skeleton code)
<https://www.instabase.com/sa3305/ieor-8100-public>
 - Login and copy the required folder to your repository (my-repo)
 - Change/write code as required
 - Submit using the submit link



Repositories

New Repository

ashipra /

 my-repo





ieor8100-spring2019 ▾



 Instabase Drive

 Notebooks

 Drives

 Databases

 Apps

 Settings

 Sharing

 Activities

 Trash

Instabase Drive

New ▾

sa3305 / ieor8100-spring2019 / fs / Instabase Drive /

Filter...




Actions ▾

▼  Labs



▼  Lab0




 Lab 0.ipynb



▶  files



▼  notebooks





Instabase Drive

Notebooks

Drives

Databases

Apps

Settings

Sharing

Activities

Trash

Instabas

sa3305

Filter...

▼ Labs

▼ Labs

► files

▼ notebooks

ieor8100-spring2019 ▼



Copy Lab0 to...



ieor8100-spring2019▼



fs / Instabase Drive /

▢ Labs



▢ files



▢ notebooks



Copy

Cancel



my-repo ▾



Instabase Drive

Notebooks

Drives

Databases

Apps

Settings

Collaborators

Activities

Trash

Instabase Drive

[ashipra](#) / [my-repo](#) / [fs](#) / Instabase Drive



☐ Lab0



☐ README.md



☐ files



☐ notebooks



☐ tutorials



New ▾

- New Folder
- New Notebook
- New File

Upload Files

README.md



Lab0

- Lab0 is a trial assignment (Jupyter notebook with skeleton code)
- Play with OpenAI gym environments, make changes to python code, plot the performance of random strategies.
- Submit using the link at the end



my-repo ▾



 Instabase Drive

 Notebooks

 Drives

 Databases

 Apps

 Settings

 Collaborators

 Activities


 Trash

Instabase Drive

New ▾

ashpra / my-repo / fs / Instabase Drive / Lab0



 Lab 0.ipynb

Preview

Open With ▸

Rename

Delete

Move

Copy

Download

Open the notebook with Jupyter

Submit by clicking on the link at the bottom (After making any changes you want. You can submit multiple times. The version will be updated every time you submit.

Secure | https://www.instabase.com/user/ashipra-nb/notebooks/ashipra/my-repo/fs/Instabase%20Drive/Lab0/Lab%200.ipynb

ib jupyter notebook

Lab 0 (autosaved)

Python 3

File Edit View Insert Cell Kernel Widgets Help

Not Trusted

Save

+

Cut

Copy

Paste

Up

Down

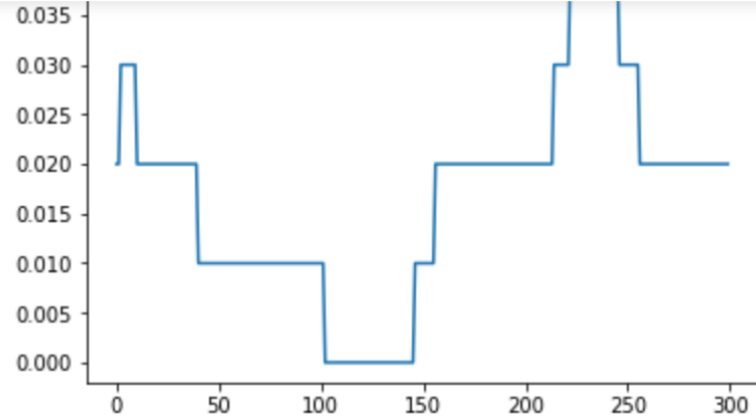
Run

Stop

Refresh

Markdown

Keyboard



x	y
0	0.020
0	0.030
10	0.020
40	0.010
100	0.000
150	0.010
160	0.020
210	0.030
220	0.035
240	0.030
250	0.020
300	0.020

In your programming assignments, you will be given a skeleton of code like above. You will be asked to make changes to the code to achieve specific tasks and then submit the notebook as your submission for the assignment. Make a submission using below. This will not be graded but will help us ensure that everything is set up correctly in your instabase account.

Submit it using the following [link](#)

TO DO

1. Sign up for Piazza
2. Signup for Instabase (after receiving token)
 - Instructions on the website <https://ieor8100.github.io/rl/cloudPlatform.html>
 - Important use your **UNI as username**, email UNI@columbia.edu
 - contact us on Piazza if you have any difficulty signing up
3. Access Lab0, submit it as trial. (will not be graded)
4. Get your computer ready for offline implementation:
Software installation instructions posted on the website
<https://ieor8100.github.io/rl/installation.html>